

7313 Take-Home Exam

In this take-home exam, I will be predicting online or offline purchases made at the Buy N Small (BnS) drugstore. As implied in the task, the selected aggregation-level will be purchases on the receipt-level. Note! With warning messages outputting the XGBoost model, the handin became 5 pages. I hope you can disregard these and consider it a 4-page hand-in.

Importing from SQL

Let us start by importing the relevant data for this classification problem. The target variable *is_online* is on the correct receipt-level aggregation-level at import. The predictor variables *purchase_date*, *age*, *gender*, and *enrollment* are also by default on the correct aggregation level as the date and customer is same for all transactions under the same receipt. The predictor variables *tot_quantity*, *tot_amount* are aggregated through the SUM() function to be expressed on receipt-level. Finally, for a creative touch I include a sub-query which gets the department code of the department that receives the most revenue from the transactions in the receipt, thus on the correct aggregation level. This department code is stored in the predictor variable *main_dept*.

```
df = dbGetQuery(con,
  "SELECT receipt_id, customer_id, purchase_date, SUM(quantity) AS tot_quantity,
    SUM(amount) AS tot_amount, MAX(is_online) AS is_online, age, gender,
    enrollment, main_dept
FROM Transactions t
  LEFT JOIN Unseen u USING (receipt_id)
  LEFT JOIN Customers c USING (customer_id)
  LEFT JOIN (
    SELECT receipt_id, dept as main_dept, max(revenue)
    FROM (
      select receipt_id, dept, sum(amount * quantity) as revenue
      FROM Transactions t
      LEFT JOIN Products p USING (item)
      GROUP BY receipt_id, dept
    ) dept_revenues
    GROUP BY receipt_id
  ) max_revenue USING (receipt_id)
WHERE id23500 = 1 ##change to your studentid
  OR is_online IS NOT NULL
GROUP BY receipt_id ")
```

Data Preparation and Feature Transformation

Before presenting the final data set we do some cleaning and feature transformation. For our classification models to work, character variables are converted to factor variables. The variable *purchase_date* is binned into *purchase_month* to remove noise and capture the long-term trend in online shopping over time. The

receipt's weekday is extracted from *purchase_date* in the variable *purchase_day*, to capture weekly cyclical patterns. A graphical motivation for these transformations of the *purchase_date* variable are shown in *Appendix A1*. Furthermore, to make the model less noisy we bin ages into generations of 5 years. This should smooth out the distribution of online receipts across age groups. A graphical motivation is found in *Appendix A2*. Due to space limitations to exact code on how these variables were transformed can be found in *Appendix A3*.

Before we proceed to the glimpse of the data, we check for missing values in the variables.

```
sapply(df, function(x) sum(is.na(x)))[sapply(df, function(x) sum(is.na(x))) != 0]
```

```
##      gender enrollment  main_dept generation
##           1           1           2           1
```

We see that there are missing values in the *generation*, *gender* and *enrollment* variables. These missing values are the result of a mismatch in *customer_id* during the joining of tables in SQL. With a full explanation in *Appendix A4*, I reformatted the faulty *customer_id* from date to numeric and found an existing customer. These correct customer details are updated to replace the missing values. For the missing values in *main_dept*, I wrote an function that imputes a *main_dept* category based on the mean *tot_amount* of each category. The mean *tot_amount* closest to the amounts in the receipts with missing *main_dept*, will determine which category is imputed.

```
# define function for imputing main_dept based on amount
impute_main_dept <- function(df, tot_amount) {
  amounts <- df %>% group_by(main_dept) %>% summarize(avg_tot_amount = mean(tot_amount)) %>%
  filter(!is.na(main_dept)) %>% arrange(desc(avg_tot_amount))
  output <- c()
  for (t in tot_amount) {
    diff_vector <- abs(t - amounts$avg_tot_amount)
    output <- c(output, as.numeric(as.character(amounts$main_dept[which.min(diff_vector)])))
  }
  return(output)
}
# call the impute_main_dept function on all receipts with NA for main_dept
df[is.na(df[, 'main_dept']), 'main_dept'] <- impute_main_dept(df, df[is.na(df[, 'main_dept']), 'tot_amount'])
```

With necessary data cleaning, feature transformation, and imputations complete, I present a glimpse of the final dataset. (Note additional feature transformation will be done for the XGBoost model).

```
glimpse(df)
```

```
## Rows: 405,411
## Columns: 9
## $ purchase_month <fct> 2019-05, 2019-08, 2019-05, 2018-10, 2018-06, 2018-06, 2~
## $ purchase_day   <fct> Friday, Tuesday, Monday, Saturday, Friday, Saturday, We~
## $ tot_quantity   <dbl> 3, 1, 2, 1, 6, 2, 3, 3, 3, 2, 7, 2, 4, 3, 6, 3, 3, 1, 1~
## $ tot_amount      <dbl> 1430, 3790, 1000, 340, 4610, 650, 2260, 2160, 1350, 197~
## $ is_online       <fct> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
## $ gender          <fct> Female, Female, Female, Female, Female, Female, Female, ~
## $ generation      <fct> 40, 55, 35, 60, 40, 35, 60, 25, 40, 60, 30, 60, 25, 40, ~
## $ enrollment      <fct> Paperform, POS, Paperform, Special, Special, Special, P~
## $ main_dept       <fct> 343, 357, 343, 349, 341, 342, 348, 339, 342, 354, 339, ~
```

Model Evaluation

I will now proceed with training, testing, and evaluating different machine learning models to predict whether receipts are made online or offline. Due to space limitations I will only present two of the two best-performing models that I explored: random forests, and XGBoost. All models will be trained and tested on the same training (75%) and testing (25%) data sets. I also declare the target variable and predictors in a formula.

```
# use a uniform distribution to split the data into training and test sets.
set.seed(7313)
dist <- runif(nrow(df))
df_train <- df[dist < 0.75,]
df_test <- df[dist >= 0.75,]

target <- "is_online"
predictors <- c("purchase_month", "purchase_day", "tot_quantity", "tot_amount", "generation", "gender",
fmla <- as.formula(paste(target, "~", paste(predictors, collapse = " + ")))
```

Random Forest

I start fitting a random forest model on the training data.

```
# Fit random forest model
online_model_rf <- ranger(fmla, df_train, num.trees = 500,
                          respect.unordered.factors = "order",
                          seed = set.seed(7313))
```

```
## Growing trees.. Progress: 60%. Estimated remaining time: 20 seconds.
```

With the model trained I now evaluate it now on the test set. I compare it to the heuristic which would be predicting all zeros for *is_online* (as it is most likely with a mean of 0.07).

```
# Make predictions on the test set
df_test$pred <- predict(online_model_rf, df_test)$predictions

# Evaluate accuracy of predictions
df_test %>% summarize(
  total = length(is_online), correct = sum(is_online == pred),
  share_correct = (100 * correct / total), incorrect = sum(is_online != pred),
  share_incorrect = (100 * incorrect / total))
```

```
##      total correct share_correct incorrect share_incorrect
## 1 100970   94633       93.72388     6337       6.276122
```

```
# Evaluate accuracy of heuristic
df_test %>% summarize(
  total = length(is_online), correct = sum(is_online == 0),
  share_correct = (100 * correct / total), incorrect = sum(is_online != 0),
  share_incorrect = (100 * incorrect / total)
)
```

```
##      total correct share_correct incorrect share_incorrect
## 1 100970   93590       92.6909      7380       7.309102
```

Comparing the model with the heuristic, we see that the model performs better than the heuristic with an accuracy of 93.72%, compared to 92.69. The difference between their accuracy and error rates is 1.03 percentage points.

XGBoost

For the XGBoost I make sure to one-hot encode all of the categorical/factor variables. The encoded training and test sets are called `df_train.treat` and `df_test.treat`, respectively. In training the model I run the XGBoost cross-validation to identify the best number of trees in the model, i.e. the number of trees that minimize the estimated out-of-sample learning error. To avoid overfitting I select the optimal number of trees from the cross-validation test sample.

```
cv <- xgb.cv(data = as.matrix(df_train.treat),
             label = as.numeric(as.character(df_train$is_online)),
             nrounds = 200, nfold = 5, objective = "binary:logistic",
             eta = 0.3, max_depth = 6, early_stopping_rounds = 10,
             verbose = 0 )

## [23:47:22] WARNING: amalgamation/./src/learner.cc:1115: Starting in XGBoost 1.3.0, the default eval
## [23:47:22] WARNING: amalgamation/./src/learner.cc:1115: Starting in XGBoost 1.3.0, the default eval
## [23:47:23] WARNING: amalgamation/./src/learner.cc:1115: Starting in XGBoost 1.3.0, the default eval
## [23:47:23] WARNING: amalgamation/./src/learner.cc:1115: Starting in XGBoost 1.3.0, the default eval
## [23:47:24] WARNING: amalgamation/./src/learner.cc:1115: Starting in XGBoost 1.3.0, the default eval

# Determine and print how many trees minimize training and test error
elog_summary <- cv$evaluation_log %>%
  summarize(ntrees.train = which.min(train_logloss_mean), # find the index of min(train_rmse_mean)
            ntrees.test = which.min(test_logloss_mean)) # find the index of min(test_rmse_mean)
(ntrees <- elog_summary['ntrees.test'][1,1])

## [1] 185
```

With the optimal number of trees determined from the cross-validation, I now run the XGBoost model and make predictions on the online/offline variable of the receipts.

```
online_model_xgb <- xgboost(data = as.matrix(df_train.treat), # training data as matrix
                             label = as.numeric(as.character(df_train$is_online)),
                             nrounds = ntrees, objective = "binary:logistic",
                             eta = 0.3, depth = 6, verbose = 0 )

## [23:52:44] WARNING: amalgamation/./src/learner.cc:576:
## Parameters: { "depth" } might not be used.
##
## This could be a false alarm, with some parameters getting used by language bindings but
## then being mistakenly passed down to XGBoost core, or some parameter actually being used
## but getting flagged wrongly here. Please open an issue if you find any such cases.
##
## [23:52:45] WARNING: amalgamation/./src/learner.cc:1115: Starting in XGBoost 1.3.0, the default eval
```


```
# Make predictions  
df_test$xgb_prob <- predict(online_model_xgb, as.matrix(df_test.treat))
```

The XGBoost model predicts a probability of receipts being online. To convert this to binary outcomes, I loop through thresholds to find the threshold which gives the highest accuracy. This is demonstrated in Appendix A5. In a similar comparison between model and heuristic (as done with random forest) we find that the model has an accuracy of 93.89%, which is 1.20 percentage points higher than the heuristic of 92.69%. As the XGBoost model performs better than the random forest I select it as the model for the unseen dataset (using Occam's razor principle of selecting simplest model given efficiency).


Appendix

A1. Splitting the purchase_date Variable

The share of receipts being online, varies over time. By splitting up the *purchase_date*, we can capture two different effects over time. When looking at the monthly development of online receipts, we see that in the months October 2017 - August 2018, there share of online receipts is negligible. After August 2018 we see an increase in online purchases. Knowing which month a purchase occurred will be very valuable in predicting the online/offline status of the purchase.




Furthermore we see that there is significant variation in the share of both purchases and online purchases between different weekdays. In the plot below, we see that Thursdays and Fridays show significantly fewer purchases as a whole, but even more so there are disproportionately fewer online purchases being made on these days. The other days exhibit quite similar levels of purchases, with a higher share of online receipts on the weekend.




A2. Binning the age Variable into 5-year Generations

The distribution of receipts (online/offline) across customer ages is shown in the graph below. In the distribution we see that there are no customers between ages 0-10, and very few older than 80.



To smoothen out the distribution, and remove noise specific to the customers in the data set, I choose to bin the customers into fixed-width bins of 5 years. The group of 5 years will be represented by the first year in the group. These groups are stored in the *generation* variable.



A3.

```
df <- df_raw

# recode variables
df <- df %>%
  mutate(
    is_online = as.factor(is_online),
    purchase_date = as.Date(purchase_date, "%Y-%m-%d"),
    purchase_month = as.factor(format(purchase_date, "%Y-%m")),
    purchase_day = factor(format(purchase_date, "%A"),
      levels = c("Monday", "Tuesday", "Wednesday", "Thursday",
        "Friday", "Saturday", "Sunday")),
    main_dept = as.factor(main_dept),
    enrollment = as.factor(enrollment),
    gender = factor(gender),
    generation = NA
  )
levels(df$gender) <- c("Female", "Male")

# group ages into generation bins of five years
generations <- seq(0, 100, 5)
for (g in generations) {
  df <- df %>%
    mutate (
      generation = ifelse(age >= g & age < (g+10), g, generation)
    )
}
df$generation <- as.factor(df$generation)

# remove the age variables from the df
df <- df %>% select(-age)
```


A4. Updating Missing Values for Customer “01-nov-00”

In the data preparation there we found missing values for age, gender and enrollment for one receipt. The reason for this was that the receipt had a *customer_id* of “01-nov-00”. If we look at a summary of the *purchase_date* distribution, we see that the equivalent date of 2000-11-01 is a severe outlier compared to the data.

```
summary(df$purchase_date)
```

```
##           Min.          1st Qu.          Median          Mean          3rd Qu.          Max.
## "2017-10-03" "2018-06-27" "2019-01-12" "2019-01-02" "2019-07-13" "2019-12-16"
```

This suggests that the *customer_id* of “01-nov-00” is not the accidental insert of a true *purchase_date*, but rather an error in variable type. If we look at the numeric equivalent of the date 2000-11-01, we find that a potential true *customer_id* could be 11262.

```
as.numeric(as.Date("01-nov-00", "%d-%b-%y"))
```

```
## [1] 11262
```

We can then check if this *customer_id* exists in the df. We present all receipts with *receipt_id* in the table below, along with the receipt with the missing values:

```
df %>%
  filter(
    (customer_id == as.numeric(as.Date("01-nov-00", "%d-%b-%y"))) | (customer_id == "01-nov-00")
  )
```

```
##           receipt_id customer_id purchase_date tot_quantity tot_amount is_online
## 1 314238412747203      11262    2019-07-23             1         500           0
## 2 379844308696576      11262    2018-11-06             2        1900           1
## 3 382324908696576    01-nov-00    2019-03-23             1        3590           1
## 4 682031087262690      11262    2018-08-30             1         590           0
## 5 701003549056052      11262    2018-10-04             1         750           0
## 6 710143849046631      11262    2019-04-25             3        1860           0
##   gender enrollment main_dept purchase_month purchase_day generation
## 1   Male         Web       354       2019-07       Tuesday        30
## 2   Male         Web       345       2018-11       Tuesday        30
## 3   <NA>         <NA>       357       2019-03       Saturday       <NA>
## 4   Male         Web       341       2018-08       Thursday        30
## 5   Male         Web       339       2018-10       Thursday        30
## 6   Male         Web       352       2019-04       Thursday        30
```

Seeing these receipts in the same table, reencoding “01-nov-00” to its numerical equivalent is justified as the *purchase_date*, *quantity*, *amount*, and *is_online* details fit in with the receipts of customer 11262. Thus we correct the missing values to a 30-35 year-old male with web enrollment.

While, this analysis took a bit more time, this process is arguable better than imputing. If we were to impute the categorical mode, then the missing values would take the values of the average customer to BnS, grouped by age, gender, and enrollment. Furthermore we would impute values into a seemingly invalid *customer_id* of “01-nov-00”. The imputed values would then be a 40-45 year-old female with paperform enrollment.

```
summary <- df %>%
  group_by(gender, generation, enrollment) %>%
  summarize (
    receipt_count = length(receipt_id)
  ) %>%
  arrange(desc(receipt_count))
```

'summarise()' has grouped output by 'gender', 'generation'. You can override using the '.groups' arg

```
summary[1:5,]
```

```
## # A tibble: 5 x 4
## # Groups:   gender, generation [5]
##   gender generation enrollment receipt_count
##   <fct>   <fct>      <fct>          <int>
## 1 Female 40      Paperform      17781
## 2 Female 30      POS           16279
## 3 Female 50      Paperform      15933
## 4 Female 45      Paperform      14999
## 5 Female 25      POS           14493
```

A5. Threshold computing for XGBoost and evaluation

```
accuracy <- c()
threshold <- seq(0.01, 1, 0.01)
for (t in threshold){
  df_test$temp_pred <- as.numeric(df_test$xgb_prob >= t)
  df_temp <- df_test %>%
  summarize(
    total = length(is_online),
    correct = sum(is_online == temp_pred),
    share_correct = (100 * correct / total),)
  accuracy <- c(accuracy, df_temp[1, 'share_correct'])
}

max_accuracy <- data.frame(threshold, accuracy) %>%
  arrange(desc(accuracy)) %>%
  filter(row_number()==1)

max_accuracy      # Evaluate accuracy of predictions    = 93.89%    (+1.20%)
```

```
##   threshold accuracy
## 1      0.47 93.89026
```

```
df_test$pred_xgb <- as.numeric(df_test$xgb_prob >= max_accuracy[1, 'threshold'])

# Evaluate accuracy of heuristic    = 92.69%
df_test %>%
  summarize(
    total = length(is_online),
    correct = sum(is_online == 0),
    share_correct = (100 * correct / total),
    incorrect = sum(is_online != 0),
    share_incorrect = (100 * incorrect / total)
  )
```

```
##   total correct share_correct incorrect share_incorrect
## 1 100970   93590      92.6909      7380      7.309102
```